

# Machine Learning and Natural Language Processing Interview Questions

March 11, 2016

- What is POS (Part of Speech) tagging? What is the simplest approach to building a POS tagger that you can imagine?
- How would you train a model that identifies whether the word "Apple" in a sentence belongs to the fruit or the company?
- Design a recommendation system.
- What is latent semantic indexing and where can it be applied?
- How would you build a system that auto corrects text that has been generated by a speech recognition system?
- What is regular grammar? Does this differ in power to a regular expression and if so, in what way?
- How would you build a POS tagger from scratch given a corpus of annotated sentences? How would you deal with unknown words?
- How would you find all the occurrences of quoted text in a news article?